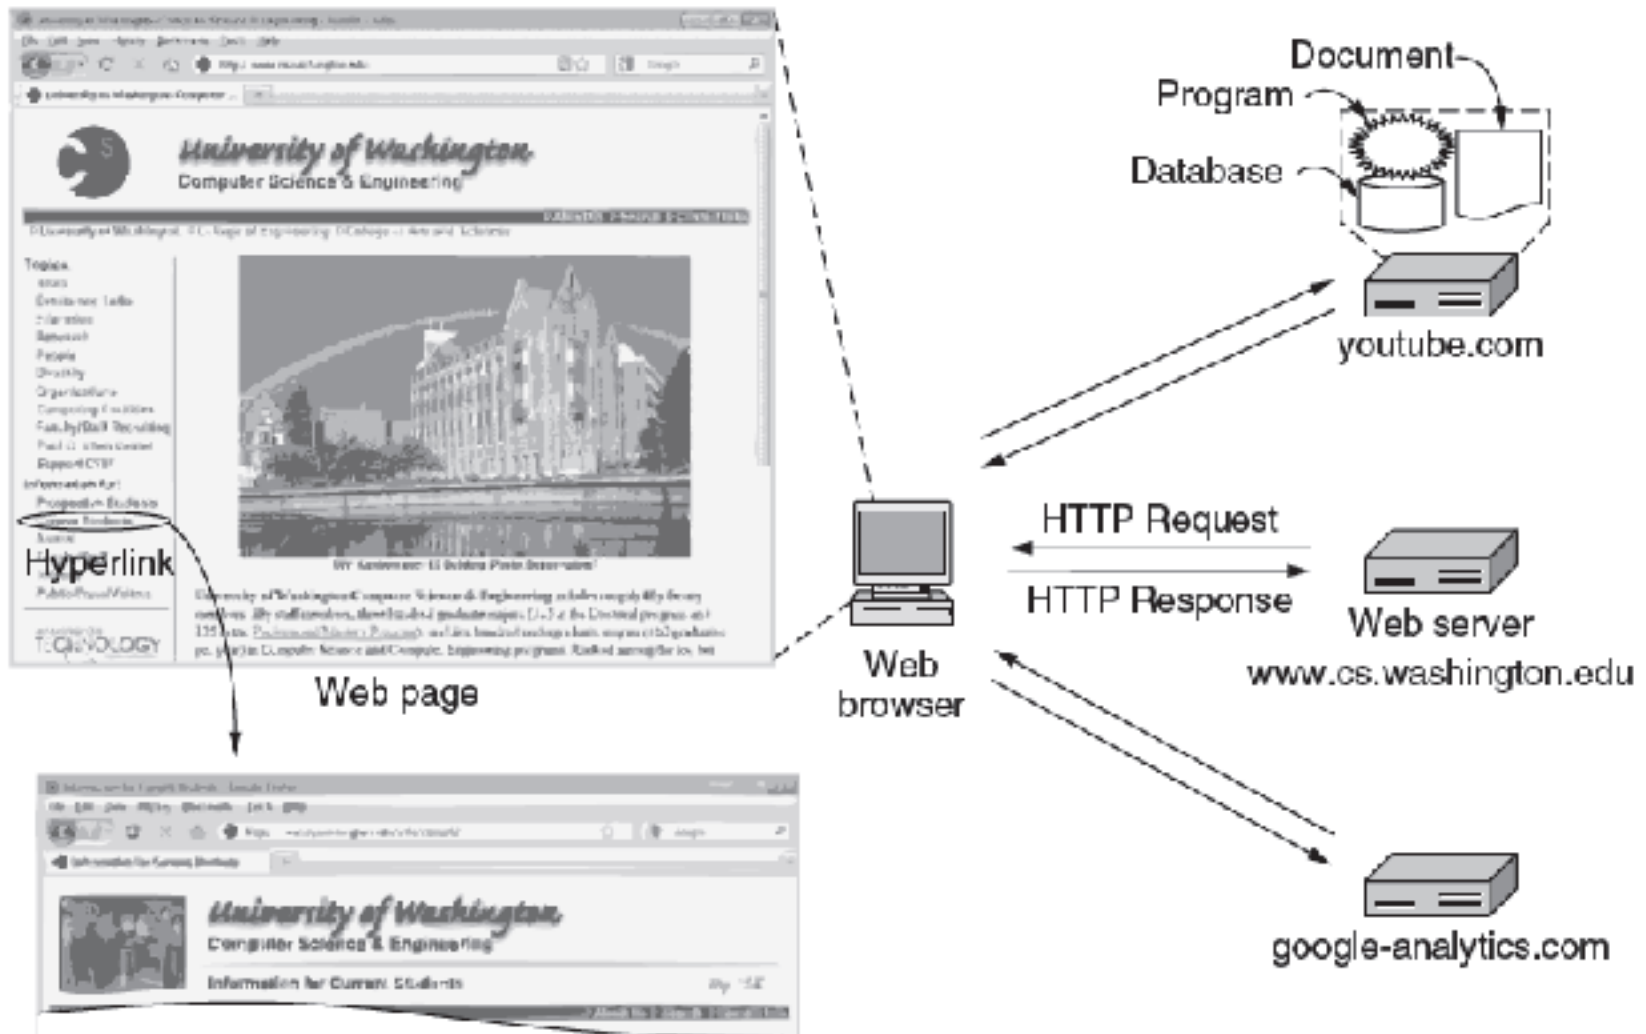


Couche application

HTTP

Computer Networks. Tanenbaum
Computer Networking. Kurose&Ross

Http: Principles



Web et HTTP

- ❖ Une *page web* contient des *objets*
- ❖ Objet : fichier HTML, images JPEG, applet, fichiers audio,...
- ❖ Une page web page consiste en un fichier de base *HTML* contenant des objets référencés
- ❖ Chaque objet est adressable par une *URL* (*Uniform Resource Locator*)

`www.someschool.edu/someDept/pic.gif`

host name

path name

URL:

- Example: <http://www.phdcomics.com/comics.php>

Protocol

Server

Page on server

Our
focus →

Name	Used for	Example
http	Hypertext (HTML)	http://www.ee.uwa.edu/~rob/
https	Hypertext with security	https://www.bank.com/accounts/
ftp	FTP	ftp://ftp.cs.vu.nl/pub/minix/README
file	Local file	file:///usr/suzanne/prog.c
mailto	Sending email	mailto:JohnUser@acm.org
rtsp	Streaming media	rtsp://youtube.com/montypython.mpg
sip	Multimedia calls	sip:eve@adversary.com
about	Browser information	about:plugins

Common URL protocols

HTTP overview

HTTP: hypertext transfer protocol

- ❖ Web's application layer protocol
- ❖ client/server model
 - **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - **server**: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview

uses TCP:

- ❖ client initiates TCP connection (creates socket) to server, port 80
- ❖ server accepts TCP connection from client
- ❖ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❖ TCP connection closed

HTTP is “stateless”

- ❖ server maintains no information about past client requests

protocols that maintain “state” are complex!

- v past history (state) must be maintained
- v if server/client crashes, their views of “state” may be inconsistent, must be reconciled

Overview

Steps a client (browser) takes to follow a hyperlink:

- Determine the protocol (HTTP)
- Ask DNS for the IP address of server
- Make a TCP connection to server
- Send request for the page; server sends it back
- Fetch other URLs as needed to display the page
- Close idle TCP connections

Steps a server takes to serve pages:

- Accept a TCP connection from client
- Get page request and map it to a resource (e.g., file name)
- Get the resource (e.g., file from disk)
- Send contents of the resource to the client.
- Release idle TCP connections

HTTP connections

non-persistent HTTP

- ❖ at most one object sent over TCP connection
 - connection then closed
- ❖ downloading multiple objects required multiple connections

persistent HTTP

- ❖ multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

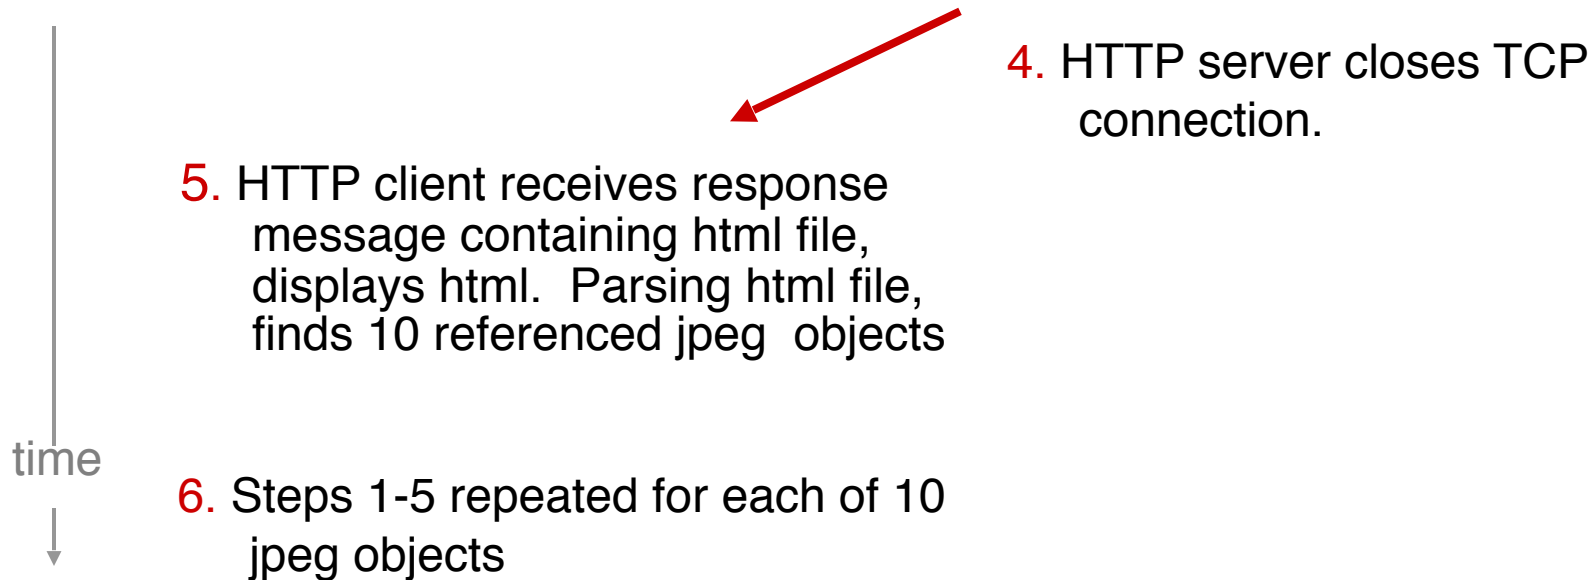
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time
↓

Carole Delporte

Non-persistent HTTP (cont.)

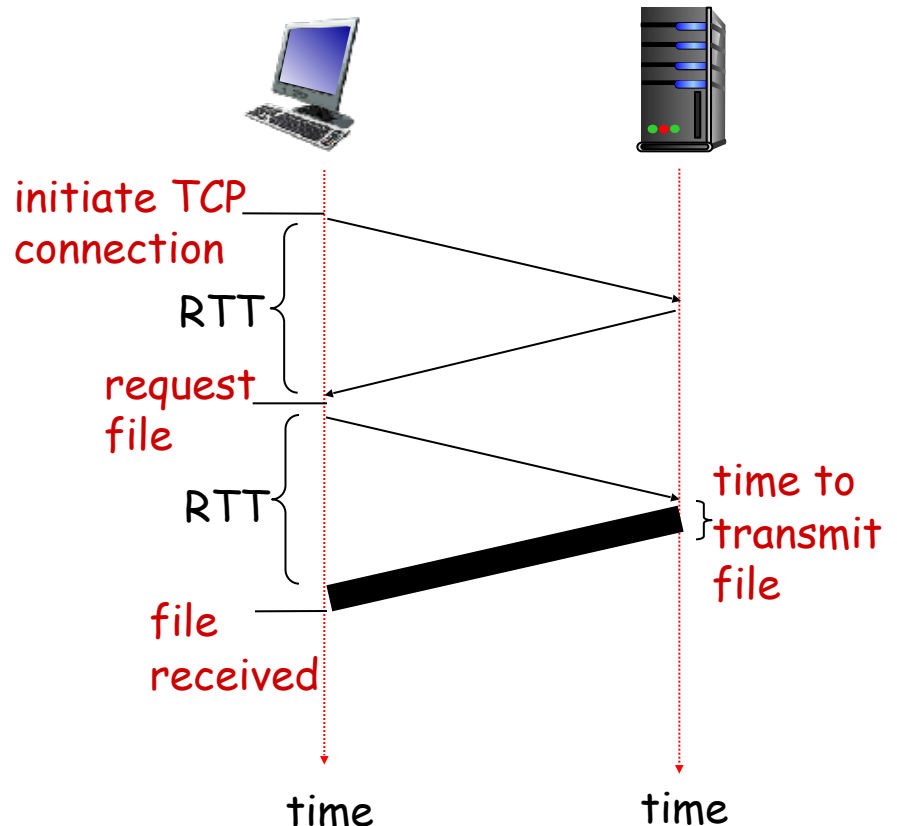


Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- ❖ one RTT to initiate TCP connection
- ❖ one RTT for HTTP request and first few bytes of HTTP response to return
- ❖ file transmission time
- ❖ non-persistent HTTP response time = $2\text{RTT} + \text{file transmission time}$



Persistent HTTP

non-persistent HTTP issues:

- ❖ requires 2 RTTs per object
- ❖ OS overhead for *each* TCP connection
- ❖ browsers often open parallel TCP connections to fetch referenced objects

persistent HTTP:

- ❖ server leaves connection open after sending response
- ❖ subsequent HTTP messages between same client/server sent over open connection
- ❖ client sends requests as soon as it encounters a referenced object
- ❖ as little as one RTT for all the referenced objects

HTTP request message

- ❖ two types of HTTP messages: *request, response*
- ❖ **HTTP request message:**
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

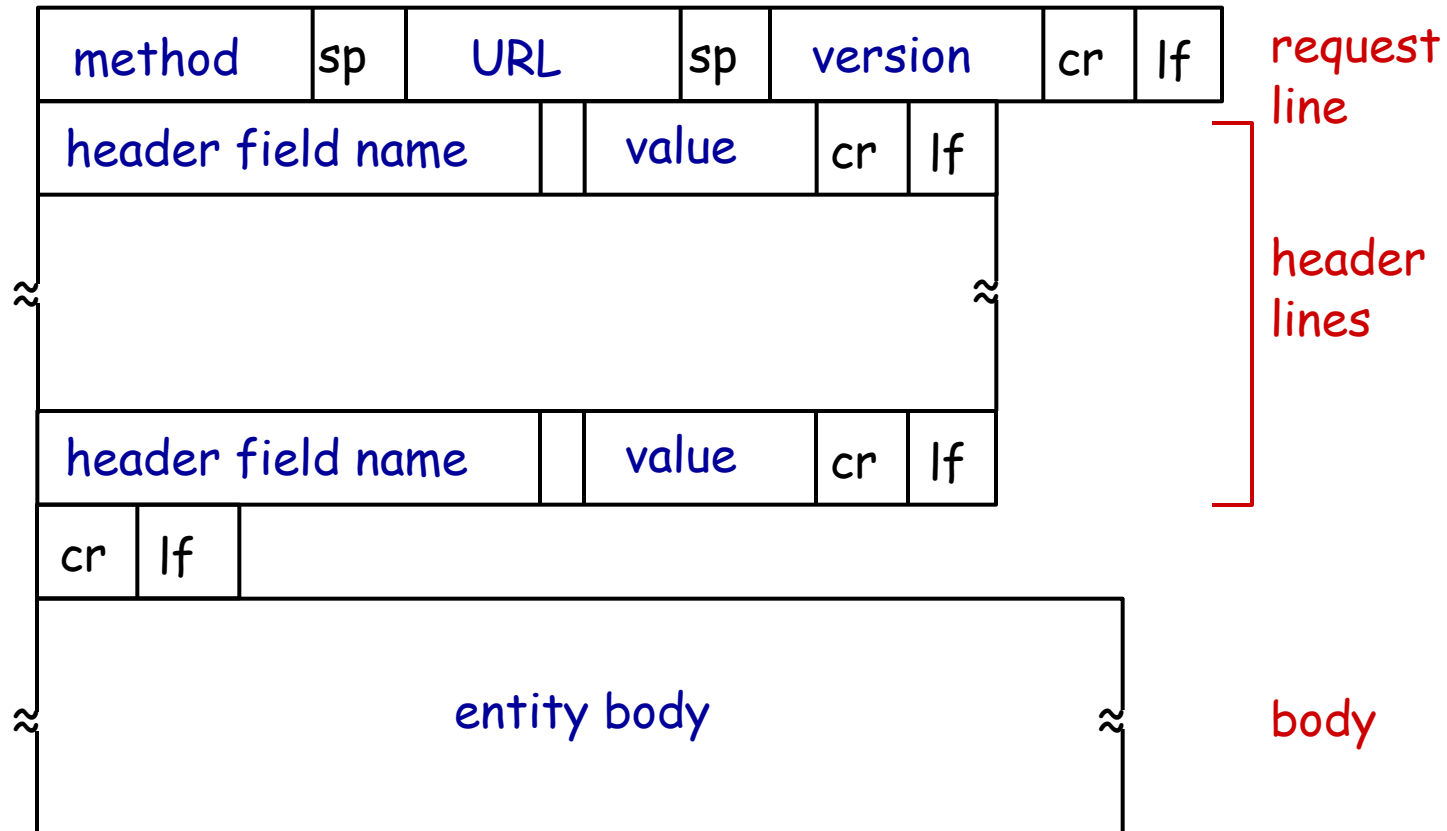
header
lines

carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

HTTP request message: general format



HTTP

Few headers:

Function	Example Headers
Browser capabilities (client → server)	User-Agent, Accept, Accept-Charset, Accept-Encoding, Accept-Language
Caching related (mixed directions)	If-Modified-Since, If-None-Match, Date, Last-Modified, Expires, Cache-Control, ETag
Browser context (client → server)	Cookie, Authorization, Host
Content delivery (server → client)	Content-Encoding, Content-Length, Content-Type, Content-Language, Set-Cookie

Uploading form input

POST method:

- ❖ web page often includes form input
- ❖ input is uploaded to server in entity body

URL method:

- ❖ uses GET method
- ❖ input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

Method types

HTTP/1.0:

- ❖ GET
- ❖ POST
- ❖ HEAD
 - asks server to leave requested object out of response

HTTP/1.1:

- ❖ GET, POST, HEAD
- ❖ PUT
 - uploads file in entity body to path specified in URL field
- ❖ DELETE
 - deletes file specified in the URL field

HTTP

Request methods.

	Method	Description
Fetch a page →	GET	Read a Web page
	HEAD	Read a Web page's header
Used to send → input data to a server program	POST	Append to a Web page
	PUT	Store a Web page
	DELETE	Remove the Web page
	TRACE	Echo the incoming request
	CONNECT	Connect through a proxy
	OPTIONS	Query options for a page

HTTP response message

status line
(protocol
status code
status phrase)

header
lines

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html;
    charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

data, e.g.,
requested
HTML file

HTTP response status codes

v status code appears in 1st line in server-to-client response message.

v some sample codes:

200 OK

- request succeeded, requested object later in this msg

301 Moved Permanently

- requested object moved, new location specified later in this msg (Location:)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

HTTP

Response codes tell the client how the request fared:

Code	Meaning	Examples
1xx	Information	100 = server agrees to handle client's request
2xx	Success	200 = request succeeded; 204 = no content present
3xx	Redirection	301 = page moved; 304 = cached page still valid
4xx	Client error	403 = forbidden page; 404 = page not found
5xx	Server error	500 = internal server error; 503 = try again later

Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

```
telnet www.irif.fr 80
```

opens TCP connection to port 80
(default HTTP server port) at www.irif.fr
anything typed in sent
to port 80 at www.irif.fr

2. type in a GET HTTP request:

```
GET /~cd/ HTTP/1.1
```

```
Host: www.irif.fr
```

by typing this in (hit carriage
return twice), you send
this minimal (but complete)
GET request to HTTP server

3. look at response message sent by HTTP server!

```
$ telnet www.liafa.univ-paris-diderot.fr 80
Trying 81.194.27.176...
Connected to www.liafa.univ-paris-diderot.fr.
Escape character is '^]'.
GET /~cd/ HTTP/1.1
Host: www.irif.fr
```

```
HTTP/1.1 302 Found
Date: Mon, 09 Oct 2017 10:21:09 GMT
Server: Apache/2.4.27 (FreeBSD) PHP/5.6.31 OpenSSL/1.0.1s-freebsd
Location: https://www.irif.fr/~cd//
Content-Length: 209
Content-Type: text/html; charset=iso-8859-1
```

```
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>302 Found</title>
</head><body>
<h1>Found</h1>
<p>The document has moved <a href="https://www.irif.fr/~cd/">here</a>.</p>
</body></html>
Connection closed by foreign host.
```

User-server state: cookies

many Web sites use cookies

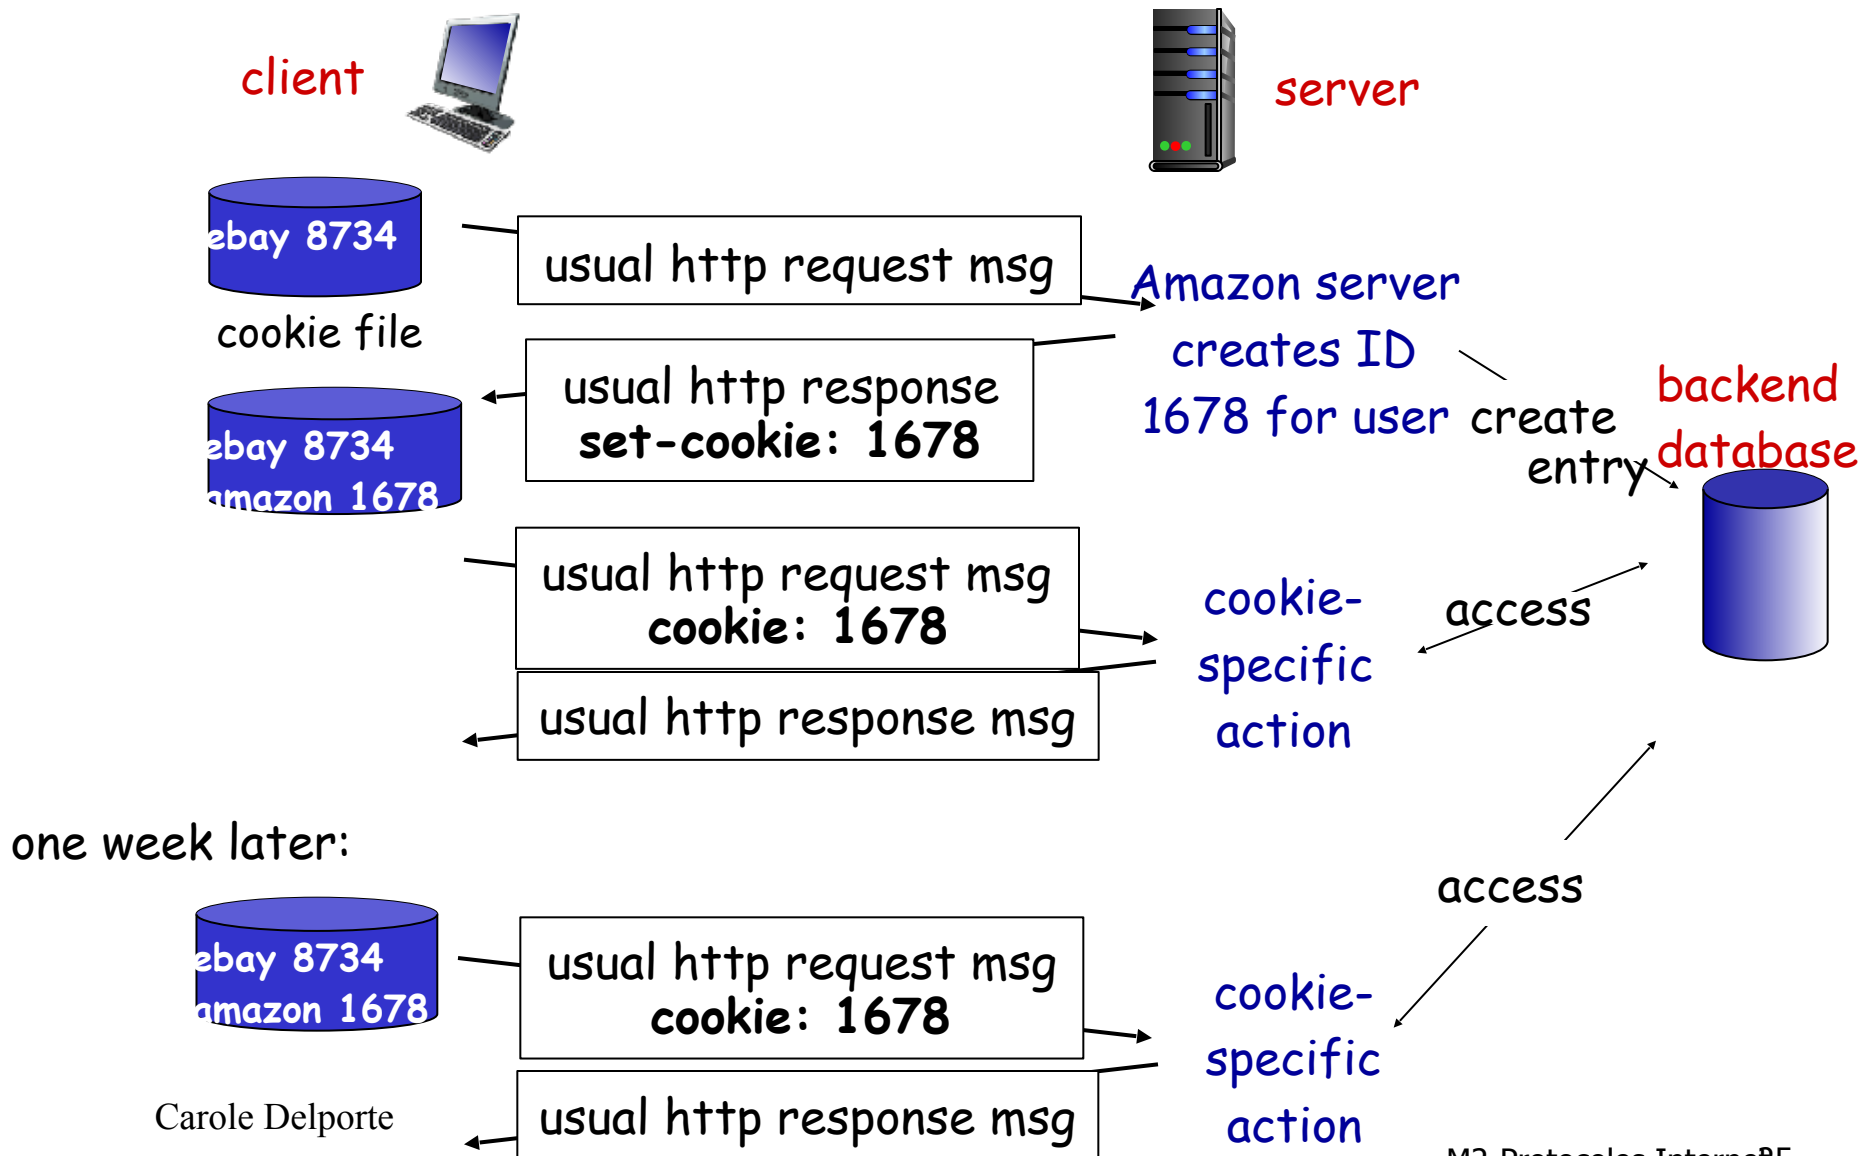
four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

example:

- ❖ Susan always access Internet from PC
- ❖ visits specific e-commerce site for first time
- ❖ when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies: keeping “state” (cont.)



Cookies (continued)

what cookies can be used for:

- ❖ authorization
- ❖ shopping carts
- ❖ recommendations
- ❖ user session state (Web e-mail)

how to keep “state”:

- v protocol endpoints: maintain state at sender/receiver over multiple transactions
- v cookies: http messages carry state


cookies and privacy.

- v cookies permit sites to learn a lot about you
- v you may supply name and e-mail to sites

Caching

To scale performance, Web servers can use:

- Caching, multiple threads, and a front end



Caching...

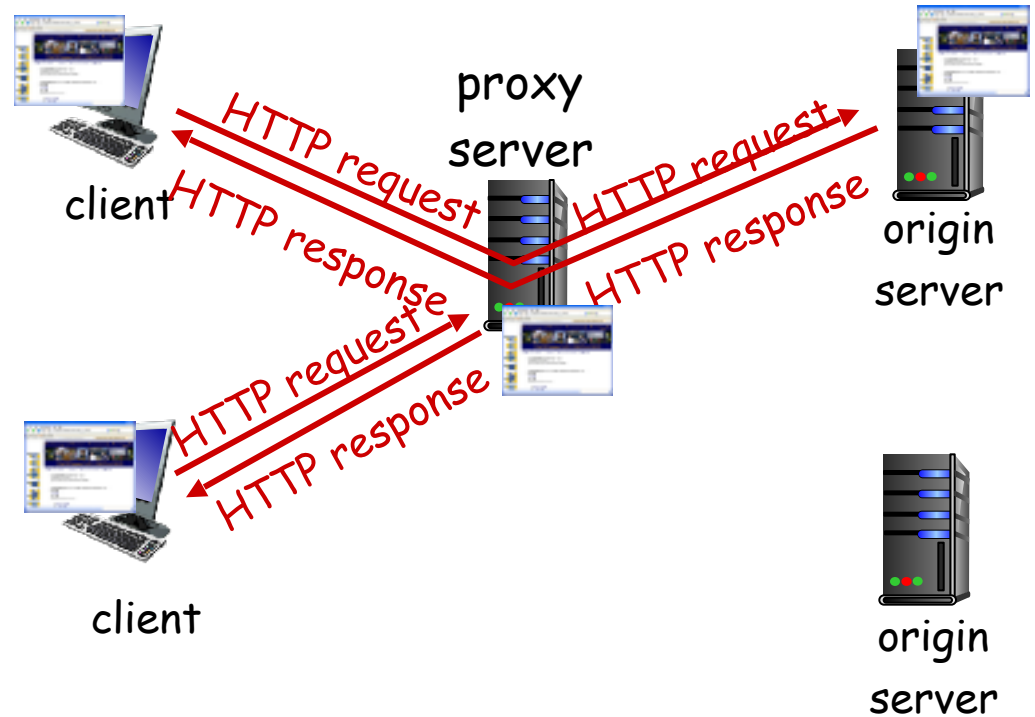
Server steps, revisited:

- Resolve name of Web page requested
- Perform access control on the Web page
- Check the cache
- Fetch requested page from disk or run program
- Determine the rest of the response
- Return the response to the client
- Make an entry in the server log

Web caches (proxy server)

goal: satisfy client request without involving origin

- ❖ **server**
- ❖ user sets browser: Web accesses via cache
- ❖ browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- ❖ cache acts as both client and server
 - server for original requesting client
 - client to origin server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- ❖ reduce response time for client request
- ❖ reduce traffic on an institution's access link
- ❖ Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)


Caching example:

assumptions:

- v avg object size: 1Mbits
- v avg request rate from browsers to origin servers: 15 request/sec
- v avg RTT from institutional router to any origin server: 2 sec
- v access link rate: 15 Mbps

consequences:

- v Traffic intensity on the LAN : 15% **problem!**
- v access link utilization = **100%**
- v total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + millisecs



Caching example: fatter access link

assumptions:


- v avg object size: 1Mbits
- v avg request rate from browsers to origin servers: 15/sec
- v RTT from institutional router to any origin server: 2 sec
- v access link rate: 15 Mbps

150 Mbps

consequences:

- v LAN utilization: 15%
- v access link utilization = 100%
- v total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + msecs

msecs



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:


- v avg object size: 1Mbits
- v avg request rate from browsers to origin servers: 15/sec
- v RTT from institutional router to any origin server: 2 sec
- v access link rate: 15 Mbps

consequences:

- v LAN utilization: 15%
- v access link utilization : 100%
- v total delay = Internet delay + LAN delay ?

How to compute link utilization, delay?

Cost: web cache (cheap!)



Caching example: install local cache

Calculating access link utilization, delay with cache:


- ❖ suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin

v access link utilization:

- § 60% of requests use access link
- § access link utilization = 60%

v total delay

- § = $0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
- § = $0.6 (2.01) + 0.4 (\sim \text{msecs})$
- § = $\sim 1.2 \text{ secs}$
- § less than with 150 Mbps link (and cheaper too!)



Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- ❖ **cache:** specify date of cached copy in HTTP request
If-modified-since:
<date>
- ❖ **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified

client



server

